Lecture 7 Data Definition Language (DDL), Views and Indexes

Instructor: Shel Finkelstein

Reference:

A First Course in Database Systems, 3rd edition, Chapter 2.3 and 8.1-8.4

Important Notices

- Lecture slides and other class information will be posted on Piazza (not Canvas).
 - Slides are posted under Resources → Lectures
 - Lecture Capture recordings are available to all students under Yuja.
 - That includes classes given over Zoom.
 - There's <u>no</u> Lecture Capture for Lab Sections.
- All Lab Sections (and Office Hours) normally meet In-Person, with rare exceptions announced on Piazza.
 - Some slides for Lab Sections have been posted on Piazza under Resources→Lab Section Notes
- Office Hours for TAs and me are posted in Syllabus and Lecture1, as well as in Piazza notice
 27; that post also now includes hours for Group Tutors.
- Some suggestions about use (and non-use) of Generative AI systems such as ChatGPT were posted on Piazza in <u>@41</u>, with specific discussion about the Movies tables and the four Avatar queries in Lecture 4.

Important Notices

- Lab2 was due Wednesday, February 7, by 11:59pm.
 - Lab2 Solution was posted on Piazza under Resources → Lab2 on Thursday,
 February 8.
 - A correction to solution for query2 was posted on Monday, Feb 12 in <u>@83</u>, and our solution was fixed.
- Lab3 will be posted on Piazza on Wednesday, Feb 14, the day of the Midterm.
 - Lab3 is due on Tuesday, February 27 by 11:59pm.
 - Lab3 has not been posted yet because:
 - a) Midterm is on Wednesday, Feb 14, and ...
 - b) There are several Lab3 topics that we haven't discussed yet.
- There will be Lab Sections before the Midterm, on topics including:
 - Lectures, Lab2 solution, last quarter's Midterm, Gradiance assignments
 - There will also be Lab Sections after the Midterm on topics including Lab3.

Important Notices

- Second Gradiance Assignment, CSE 180 Winter 2024 #2, was due on Tuesday, February 6, by 11:59pm.
- Third Gradiance Assignment, CSE 180 Winter 2024 #3, was assigned on Tuesday, February 6.
 - It is due on Thursday, February 15, by 11:59pm, the day after the Midterm.
 - But it's a good idea to complete this Gradiance Assignment before the Midterm.
 - The day of the week on which Gradiance Assignments are due will vary.
 - There are 9 questions in this Gradiance Assignment.
 - Some of them may be difficult.
 - Website and Class Token for Gradiance are in the Syllabus and first Lecture.
 - If you forget your Gradiance password, you can recover it by providing your Gradiance login and email.
 - You may take each Gradiance Assignment as many times as you like.
 - Questions stay the same, but answers change.
 - Your Gradiance Assignment score is the score on your <u>Final</u> Gradiance submission <u>before</u> the deadline.
 - Gradiance does not allow late submissions, so no extensions are possible.

Very Important Notice

- The Winter 2024 CSE 180 Final is on Wednesday, March 20 from 8:00am 11:00am in our usual classroom.
 - The Syllabus and First Lecture originally said that it was on Tuesday,
 March 19.
 - That's been fixed. Please adjust your calendars accordingly!

Important Notices: Midterm

CSE 180 Midterm is on Wednesday, February 14, and it will be given in-person, except for students who receive Remote Exam permission due to serious illnesses.

- Usual class time, 65 minutes (extended for DRC students).
 - Hope that all DRC students have submitted their forms for accommodations...
- No early/late Exams. No make-up Exams. No devices.
- You may bring one double-sided 8.5" x 11" sheet to the Midterm, with anything that you want written or printed on it that you can read unassisted) ...
 - (Okay, you may bring two sheets with writing on only <u>one side</u> of each sheet.)
 - But you must not use any other material or receive any help during the
 Midterm, whether you're in the classroom, remote, or in a DRC room
 - As the Syllabus emphasize, Academic Integrity violations have serious consequences!
- Write your answers on the exam itself, readably, using either ink or a #2 pencil.
- We might assign seats to you as you enter the room.

Important Notices: Midterm (continued)

- If you need to take the Midterm Remotely due to serious illness, please send me an email whose subject is "Taking CSE 180 Midterm Remotely" justifying that request.
 - Send that email <u>before 8:00pm</u> on Tuesday, February 13, the day before the Midterm. I'll send you instructions before 10:30am on Wednesday, February 14.
 - Only students whose requests have been <u>approved</u> may take the Midterm Remotely.
- The CSE 180 Midterm from Winter 2023 was posted under Resources → Exams on Piazza on Wednesday,
 February 1.
 - We haven't covered all the material in that previous Midterm yet.
 - Also, last year's Midterm used that quarter's database schema; your schema is different.
 - Solution to Winter 2023 Midterm <u>has also been posted on Piazza</u> ... but I strongly suggest that you take it yourself first, rather than just reading the solution.
- All past Lectures including the first half of this one, Lecture 7 (Views) will be included on the Midterm.
 - The second half of Lecture 7 (Indexes) won't be on the Midterm, even if we discuss it in class.
- At the end of the Midterm, you'll hand in your Midterm paper, and you'll show your UCSC ID.
 - Do not hand in your 8.5" x 11" sheet.
- There <u>will</u> be Lab Sections and Tutoring during the rest of the week after the Midterm, on topics including:
 - Lectures, Lab2 solution, Gradiance, Lab3.
 - Lab3 will be posted on Piazza on February 14, after the Midterm.

SQL Language

- Data Manipulation Language (DML)
 - Access and modify data
 - SELECT, INSERT, DELETE, UPDATE
- Data Definition Language (DDL)
 - Modify structure of data
 - CREATE, DROP, ALTER
- Data Control Language (DCL)
 - Control access to the data (security)
 - GRANT, REVOKE
- Databases also have Utilities, such as Backup/Restore and Load.
 - Syntax not specified in the SQL standard

CREATE TABLE

```
create table MovieStar (
starName CHAR(30),
address VARCHAR(255) DEFAULT 'Hollywood',
gender CHAR(1),
birthdate DATE NOT NULL DEFAULT '2001-12-30',
PRIMARY KEY (starName)
);
```

- PRIMARY KEY
- DEFAULT
- NOT NULL

DROP TABLE

Dropping a table:

DROP TABLE MovieStar;

- Don't assume that rolling back transaction will bring back the table!
 - Interaction of DDL and transactions may depend on implementation.

ALTER TABLE

- Adding a column to a table:
 - ALTER TABLE MovieStar ADD phone CHAR(16) DEFAULT 'unlisted';
- Dropping a column from a table:
 - ALTER TABLE MovieStar DROP birthdate;
 - In some systems:
 ALTER TABLE MovieStar DROP COLUMN birthdate;
 - In some SQL systems, dropping a column isn't allowed.
- Changing the type of a column:
 - Some implementations let you change type of column in limited ways.

What Can You CREATE/DROP in SQL DDL?

- TABLE
- VIEW
- INDEX
- ASSERTION
- TRIGGER
- SCHEMA
- PROCEDURE/FUNCTION/TYPE
 - SQL2003 standard, but there are significant variations in implementations in different systems

• ...

VIEWS: Motivation for Views

 Views help with logical data independence, allowing you to retrieve data if it matches the description in the view.

```
CREATE VIEW <view-name> AS <view-definition>;
```

CREATE VIEW ParamountMovies AS

SELECT movieTitle, movieYear

FROM Movies

WHERE studioName = 'Paramount';

- You may now ask queries on ParamountMovies as if it were a table:
 SELECT movieTitle FROM ParamountMovies WHERE movieYear=1976;
 - Composition in SQL is powerful: Tables, Queries, Views

Some Advantages of Views

- Short-hand/encapsulation: You can treat a view as a table (for queries),
 which allows you to define a short-hand for a concept that might involved
 a complicated query.
- Re-use: You can re-use a view as often as you like; other people who can
 access the view may also re-use it.
- Authorization: People may be granted access to a view, even though they don't have access to the underlying tables.
 - They may not even know that the view isn't a table.
 - Even if they know that it's a view, they may not know which tables underlie it.
- Logical data independence: Even if the tables underlying a view change, the view may still be used in queries (or applications), without re-writing the queries (or applications).

More Views

Movies (<u>movieTitle</u>, <u>movieYear</u>, length, genre, studioName, producerC#) MovieExec (execName, address, <u>cert#</u>, netWorth)

CREATE VIEW MovieProd AS

SELECT m.movieTitle, e.execName, m.length

FROM Movies m, MovieExec e

WHERE m.producerC# = e.cert#;

SELECT DISTINCT length

FROM MovieProd

WHERE execName = 'George Lucas';

Renaming Attributes in CREATE VIEW

```
Movies (<u>movieTitle</u>, <u>movieYear</u>, length, genre, studioName, producerC#)
MovieExec (execName, address, <u>cert#</u>, netWorth)
```

```
CREATE VIEW MovieProd(theTitle, theProducer, theLength) AS

SELECT m.movieTitle, e.execName, m.length

FROM Movies m, MovieExec e

WHERE m.producerC# = e.cert#;
```

SELECT DISTINCT theLength
FROM MovieProd
WHERE theProducer = 'George Lucas';

What is a View?

- A view can include any SQL SELECT statement
 - Including UNION, Aggregates, GROUP BY, HAVING,
 ORDER BY, etc.
- A view is <u>not</u> stored as a table
 - The tables underlying the view are stored in the database,
 but only the <u>description</u> of the view is in the database.
- But a view can be used in many (not all) of the same ways as tables
 - Views can be queried.
 - Views can be defined on views, as well as on tables!

Queries on Views and Tables

```
CREATE VIEW Paramount Movies AS
    SELECT movieTitle , movieYear
    FROM Movies
    WHERE studioName = 'Paramount';
SELECT DISTINCT s.starName
FROM ParamountMovies p , StarsIn s
WHERE p.movieTitle = s.movieTitle AND p.movieYear = s.movieYear ;
CREATE VIEW ParamountStars AS
   SELECT DISTINCT s.starName
   FROM ParamountMovies p, StarsIn s
   WHERE p.movieTitle = s.movieTitle AND p.movieYear = s.movieYear ;
```

DROP VIEW

CREATE VIEW ParamountMovies AS

SELECT movieTitle, movieYear

FROM Movies

WHERE studioName = 'Paramount';

DROP View ParamountMovies;

- What happens if you execute the following after dropping that view?
 - SELECT * FROM ParamountMovies;
 - SELECT * FROM Movies;

Materialized Views

- Normally, only the definition of a View is stored in the database; the contents of that view are not stored.
 - That's what we'll assume in this class!
- Some database systems support Materialized Views, where users may request that the tuples in a view be physically stored in the database.
 - However
- If the relations used in the query that defines the view are updated,
 then the Materialized View contents becomes out-of-date.
 - Maintaining the Materialized View would require changing the view contents whenever its underlying relations are updated.
 - This can be expensive, so it's unusual to maintain Materialized Views.
 - Moreover, some views are not updatable ... as we'll soon see.

View Updates

- Some modification operations on views work, but others do not, generally failing either because:
 - Constraint on underlying table would be violated, or
 - The effects of the View modification is not well-defined on the underlying tables.
- This is a complex topic, which we'll only discuss briefly.
 - See Textbook Section 8.2 for more info.

View Update Problems

Movies(<u>movieTitle</u>, <u>movieYear</u>, length, genre, studioName, producerC#)

CREATE VIEW ParamountMovies AS

SELECT movieTitle, movieYear

FROM Movies

WHERE studioName = 'Paramount';

INSERT INTO ParamountMovies VALUES ('StarTrek', 1979);

The INSERT will fail if some other column of Movies (besides movieTitle and movieYear) doesn't have a default, and that column also doesn't allow NULL values.

View Update Problems (continued)

Ambiguous View Update example with Employees and Departments

<< We'll draw this on the board >>

INDEXES: Motivation for Indexes

Searching an entire table may take a long time:

SELECT *

FROM Movies

WHERE studioName = 'Disney' AND movieYear = 1990;

If there were 100 Million movies, searching them might take a while. An index (e.g., a B-Tree) would allow faster access to matching movies.

If a table is updated, Indexes on that table are immediately <u>automatically</u> updated within the same transaction.

- Which indexes do you need to change on INSERT and DELETE?
- What about UPDATE?

CREATE INDEX

```
SELECT *
FROM Movies
WHERE studioName = 'Disney' AND movieYear = 1990;
How much would each of these indexes help?
   CREATE INDEX YearIndex ON Movies(movieYear);
   CREATE INDEX StudioIndex ON Movies(studioName);
   CREATE INDEX YSIndex ON Movies(movieYear, studioName);
   CREATE INDEX SYIndex ON Movies(studioName, movieYear);
```

How much would each of the indexes help if the WHERE clause was just movieYear = 1990?

Indexes and Ordering

```
SELECT *
FROM Movies
WHERE studioName = 'Disney' AND movieYear < 1990;
How much would each of these indexes help?
   CREATE INDEX YearIndex ON Movies(movieYear);
   CREATE INDEX StudioIndex ON Movies(studioName);
   CREATE INDEX YSIndex ON Movies(movieYear, studioName);
   CREATE INDEX SYIndex ON Movies(studioName, movieYear);
```

How much would each of the indexes help if the WHERE clause was just movieYear < 1990?

Indexes and Physical Independence

- SQL statements can be executed regardless of which indexes (if any) exist in the database.
 - Applications don't have to be modified when indexes are created or dropped!
- What gets impacted when indexes are created or dropped?
 - Performance of SQL statements
 - Some may run faster, some may run slower

Disadvantages of Indexes?

- Why not put indexes on every attribute, or even on every combination of attributes that you might query on?
 - Huge number of indexes
 - Space for indexes
 - Cache impact of searching indexes
 - Update time for indexes when table is modified

Index Design

- Most Database Administrators (DBAs) pick a set of indexes that work well on expected workload, and there are tools that help pick good indexes
 - But workloads change, so choice of indexes may need to change
 - DROP INDEX YearIndex;
- Keys are indexed (automatically in many database systems) to:
 - Help maintain uniqueness (primary key, unique)
 - Check Foreign Key references to Primary Keys (Referential Integrity)

Index Utilization

- SQL statements don't specify use of indexes, so they don't have to be modified when you change what's indexed!
 - Database Optimizer tries to figure out "the best"/"a good" way to execute each SQL query.
 - All the tuples in a Relation can be scanned directly, without using indexes, so indexes aren't necessary ... except for performance.
 - Some systems have ways that you can tell the Optimizer what to do.
 This has advantages and disadvantages. (What are they?)
- Many SQL systems (including PostgreSQL) have an EXPLAIN PLAN statement, so that you can see what plan the optimizer chooses for a SQL statement.