

Lecture 8

Constraints and Triggers

Instructor: Shel Finkelstein

Reference:

*A First Course in Database Systems,
3rd edition, Chapter 7 (but not section 7.4)*

Important Notices

- Lecture slides and other class information will be posted on [Piazza](#) (not Canvas).
 - Slides are posted under Resources→Lectures
 - Lecture Capture recordings are available to all students under Yuja.
 - That includes classes given over Zoom.
 - There's no Lecture Capture for Lab Sections.
- All Lab Sections (and Office Hours) normally meet In-Person, with rare exceptions announced on Piazza.
 - Some slides for Lab Sections have been posted on Piazza under Resources→Lab Section Notes
- Office Hours for TAs and me are posted in Syllabus and Lecture1, as well as in Piazza notice [@7](#); that post also now includes hours for Group Tutors.
- Some suggestions about use (and non-use) of Generative AI systems such as ChatGPT were posted on Piazza in [@41](#), with specific discussion about the Movies tables and the four Avatar queries in Lecture 4.
- Monday, February 19 is a holiday, Presidents day.
 - There are no Lectures, Lab Sections, Office Hours or Tutoring on that day.
 - Please attend a different Lab Section if you have a Monday Lab Section.

Important Notices

- CSE 180 Midterm was on **Wednesday, February 14**.
 - Midterm and Midterm answers were posted on Piazza under Resources→Exams on Wednesday, February 14.
 - Answers will be/were discussed at the beginning of class on Friday, February 16.
 - We'll finish grading and post grades as soon as possible.
- The Winter 2024 CSE 180 Final is on **Wednesday, March 20** from 8:00am – 11:00am in our usual classroom.
- Third Gradiance Assignment, CSE 180 Winter 2024 #3, was assigned on Tuesday, February 6.
 - It was due on **Thursday, February 15, by 11:59pm**, the day after the Midterm.
 - There will be an announcement when we assign Gradiance #4, which will be after we complete much of Lecture 9, Relational Algebra.
- Lab2 grades were unmuted on Tuesday, February 13.
 - Please contact your TA and then me if you have questions about your grade by **Monday, February 19**.
 - Never contact the Reader who graded your Lab Assignment.

Important Notices

- Lab3 was posted on Piazza on Wednesday, February 14.
 - Lab3 is due on **Tuesday, February 27**.
 - Lab3 has many parts, and some parts of Lab3 are difficult.
 - We haven't discussed all of the Lab3 topics yet.
 - You'll have enough information to do Lab3 after we complete Lecture 8 (this lecture), probably on Friday, February 16.
 - You won't need to use Triggers for Lab3.
 - You must use the Lab3 create file and load_lab3.sql file to do Lab3.
 - This original load data must be reloaded at multiple stages of Lab3.
 - The create_lab3.sql file was edited on Thursday, February 15 at about noon.
 - If you downloaded it before that, please download it again.

A Word to the Unwise

- This is a tough class for some students since it involves a combination of theory and practice.
 - The second half of CSE 180 is much harder than the first half of the course.
- Students who regularly attend Lectures and Lab Sections (and Office Hours and Tutoring) often do well; students who don't regularly attend often do poorly.
 - We don't take attendance; you're responsible for your own choices.
- After the course ends, your course grade will be determined by your scores on Exams, Labs and Gradiance, as described on the “Course Evaluation” and “Grading” Slides in the Syllabus.
 - You won't be able to do any additional work to improve your grade.

Constraints and Triggers

- A *constraint* is a relationship among data elements that the DBMS is required to enforce.
 - *Example*: key constraints.
- *Triggers/Rules* are only executed when a specified condition occurs, e.g., insertion of a tuple.
 - Easier to implement than complex constraints.

Kinds of Constraints

- Primary Key/Unique constraints
- Foreign Key, or referential-integrity constraints
- Attribute-based constraints
 - Constrain values of a particular attribute
- Tuple-based constraints
 - Relationship among components of tuple
- Assertions
 - Any SQL boolean expression (not implemented in most relational DBMS, not discussed in this lecture)

Review: Single-Attribute Keys

- Place PRIMARY KEY or UNIQUE after the type in the declaration of the attribute.
- Example:

```
CREATE TABLE Beers (  
    name    CHAR(20)  UNIQUE,  
    manf    CHAR(20)  
);
```


Review: Multi-Attribute Key

- The bar and beer together are the key for Sells:

```
CREATE TABLE Sells (  
    bar    CHAR(20) ,  
    beer   VARCHAR(20) ,  
    price  REAL,  
    PRIMARY KEY (bar, beer)  
);
```

Review: NULL

```
CREATE TABLE Sells (  
    bar          CHAR(20) ,  
    beer         VARCHAR(20) ,  
    price REAL NOT NULL ,  
    PRIMARY KEY (bar, beer)  
);
```

If the CREATE statement didn't include NOT NULL for price:

```
ALTER TABLE Sells ALTER COLUMN price SET NOT NULL;
```

```
ALTER TABLE Sells ALTER COLUMN price DROP NOT NULL;
```

Foreign Keys

- Values appearing in attributes of one relation must also appear together in specific attributes of another relation.
- **Example:**
 - In **Sells(bar, beer, price)**, we might expect that a **beer** value also appears in Beers.name (the name column of the Beers table, the primary key for that table).
- Like a link/pointer, but based on **value**.

CREATE TABLE for MovieExec

```
CREATE TABLE MovieExec (  
    execName      CHAR(30),  
    address       VARCHAR(255),  
    cert#         INT PRIMARY KEY,  
    netWorth      INTEGER  
);
```

MovieExec(execName, address, cert#, netWorth)

CREATE TABLE for Studio

```
CREATE TABLE Studio (  
    studioName      CHAR(30) PRIMARY KEY,  
    address         VARCHAR(255),  
    presC#          INT,  
    FOREIGN KEY (presC#) REFERENCES MovieExec(cert#)  
);
```

Studio(studioName, address, presC#)

- Every presC# in Studio must be the Primary Key (cert#) of a movie exec in the MovieExec table.
 - But for each movie exec in the MovieExec table, there can be n studios in the Studios table for which that movie exec is president, where n can be 0, 1, 2, ..., 17, ...

CREATE TABLE for Movies

```
CREATE TABLE Movies (  
    movieTitle    CHAR(100),  
    movieYear     INT,  
    length        INT,  
    genre         BOOLEAN,  
    studioName    CHAR(30),  
    producerC#    INT,  
    PRIMARY KEY (movieTitle, movieYear),  
    FOREIGN KEY (studioName) REFERENCES Studio,  
    FOREIGN KEY (producerC#) REFERENCES MovieExec(cert#)  
);
```

Didn't have to say:
FOREIGN KEY (studioName)
REFERENCES Studio (**studioName**),
because referencing attributes and
referenced attributes are the same.

Movies(movieTitle, movieYear, length, genre, studioName, producerC#)

- Every **studioName** in Movies must be the Primary Key (**studioName**) of a studio in the Studio table.
 - But for each **studioName** in the Studio table, there can be **n** movies in the Movies table which have that **studioName**, where **n** can be 0, 1, 2, ..., 17, ...
- Every **producerC#** in Movies must be the Primary Key (**cert#**) of a movie exec in the MovieExec table.
 - But for each movie exec in the MovieExec table, there can be **n** movies in the Movies table for which that movie exec is producer, where **n** can be 0, 1, 2, ..., 17, ...

CREATE TABLE for MovieStar

```
CREATE TABLE MovieStar (  
    starName      CHAR(30) PRIMARY KEY,  
    address       VARCHAR(255),  
    gender        CHAR(1),  
    birthdate     DATE  
);
```

MovieStar(starName, address, gender, birthdate)

CREATE TABLE for StarsIn

```
CREATE TABLE StarsIn (  
    movieTitle CHAR(100),  
    movieYear INT,  
    starName      CHAR(30)  
    PRIMARY KEY (movieTitle, movieYear, starName),  
    FOREIGN KEY (movieTitle, movieYear) REFERENCES Movies,  
    FOREIGN KEY (starName) REFERENCES MovieStar  
);
```

StarsIn(movieTitle, movieYear, starName)

- Every (movieTitle, movieYear) in StarsIn must be the Primary Key of a movie in the Movies table.
 - But for each (movieTitle, movieYear) in the Movies table, there can be n Movies in the StarsIn table which have that (movieTitle, movieYear) , where n can be 0, 1, 2, ..., 17, ...
- Every starName in StarsIn must be the Primary Key of a movie star in the MovieStars table.
 - But for each starName in the MovieStar table, there can be n Movies in the StarsIn table which have that starName, where n can be 0, 1, 2, ..., 17, ...

Expressing Foreign Keys

- Use keyword REFERENCES, either:
 1. After an attribute (for one-attribute keys)
 2. As an element of the schema:
FOREIGN KEY (<list of attributes>
REFERENCES <relation> [(<list of attributes>)]
- Referenced attributes must be declared as either PRIMARY KEY or UNIQUE.
 - (Why?)

Example: With Attribute

```
CREATE TABLE Beers (  
    name    CHAR(20) PRIMARY KEY,  
    manf    CHAR(20)  
);
```

```
CREATE TABLE Sells (  
    bar CHAR(20),  
    beer CHAR(20) REFERENCES Beers(name),  
    price REAL,  
    PRIMARY KEY(bar, beer)  
);
```

Example: As a Schema Element

```
CREATE TABLE Beers (  
    name    CHAR(20) PRIMARY KEY,  
    manf    CHAR(20)  
);
```

```
CREATE TABLE Sells (  
    bar      CHAR(20),  
    beer     CHAR(20),  
    price    REAL,  
    FOREIGN KEY(beer)  
        REFERENCES Beers(name)  
);
```

Example: As a Schema Element, With a **Name for Constraint**

```
CREATE TABLE Beers (  
    name    CHAR(20) PRIMARY KEY,  
    manf    CHAR(20)  
);
```

```
CREATE TABLE Sells (  
    bar      CHAR(20),  
    beer     CHAR(20),  
    price    REAL,  
    CONSTRAINT BeerNameFK FOREIGN KEY(beer)  
        REFERENCES Beers(name)  
);
```

Adding/Dropping Foreign Key Constraint

```
CREATE TABLE Beers (  
    name CHAR(20) PRIMARY KEY,  
    manf CHAR(20)  
);
```

If we didn't include the BeerNameFK constraint in the CREATE statement ...

```
CREATE TABLE Sells (  
    bar CHAR(20),  
    beer CHAR(20),  
    price REAL,  
    PRIMARY KEY(bar, beer)  
);
```

... then we can add it, with a name, optionally...

```
ALTER TABLE Sells  
    ADD CONSTRAINT BeerNameFK FOREIGN KEY(beer)  
        REFERENCES Beers(name);
```

... and any named constraint can also be dropped.

```
ALTER TABLE Sells  
    DROP CONSTRAINT BeerNameFK;
```

Enforcing Foreign Key Constraints (Referential Integrity, RI)

If there is a Foreign Key constraint from *referring relation R (Child)* to *referenced relation S (Parent)*, then violations can occur in two ways:

1. An INSERT or UPDATE to *R* introduces values that are not found in *S*, or
2. A DELETE or UPDATE to *S* causes some tuples of *R* to “dangle”, referencing a value that no longer exists

Actions Taken --- (1)

Example: Suppose $R = \text{Sells}$, $S = \text{Beers}$.

– That is, Sells refers to Beers

- An INSERT or UPDATE to **Sells** that introduces a non-existent beer must be rejected.
- A DELETE or UPDATE to **Beers** that removes a name that appears as a beer in some tuples of **Sells** can be handled in one of three ways (next slide).

Actions Taken --- (2)

1. **RESTRICT**: Reject the modification. This is the Default.
2. **CASCADE**: Make the same changes in Sells.
 - If DELETE of a beer in Beers that has Beers.name:
Delete the Sells tuples whose Sells.beer attribute corresponds to Beers.name
 - If UPDATE of the name of a beer in Beers:
Change the Sells.beer attribute of the tuples in Sells that corresponded to that old value of Beers.name ...
 - ... so that Sells.beer has the same new value as Beers.name
3. **SET NULL**: Change Sells.beer to NULL because there's no longer a corresponding tuple in Beers with that Beers.name

Example: Cascade

- Upon DELETE of the Bud tuple from Beers:
 - Delete all tuples from Sells that have beer = 'Bud'
- Upon UPDATE of the Bud tuple, changing 'Bud' to 'Budweiser':
 - Change all Sells tuples that have beer = 'Bud' to have beer = 'Budweiser'

Example: Set NULL

- Upon DELETE of the Bud tuple from Beers:
 - Change all tuples of Sells that have beer = 'Bud' so that their beer value becomes NULL.
- Upon UPDATE of the Bud tuple, changing 'Bud' to 'Budweiser':
 - Also change all tuples of Sells that have beer = 'Bud' so that their beer value becomes NULL.
- If Sells.beer can be NULL, then the relationship of Beers.name (“parent”) to Sells.beer (“child”) is 0/1:n, not 1:n.
 - That is, every “child” Foreign Key has a parent Primary Key, except for the children whose Foreign Key value is NULL.
 - Yes, you can set attribute values to NULL.

Choosing a Referential Integrity Policy

- When we declare a Foreign Key, we may choose policies SET NULL or CASCADE independently for DELETE and UPDATE, or stay with the Default (RESTRICT).
- Policy appears in the Foreign Key declaration with:
ON [UPDATE, DELETE][SET NULL, CASCADE]
- Two such clauses may appear, one for UPDATE and one for DELETE
- If no policy is specified, then the Default (Reject) is used.
 - Specifying NO ACTION or RESTRICT explicitly specifies this default policy.
 - NO ACTION and RESTRICT are similar, but there is a subtle difference between them as to when constraint is checked.
 - Don't specify “Reject” as the policy.

Example: Setting Policy

```
CREATE TABLE Sells (  
    bar      CHAR(20),  
    beer     CHAR(20),  
    price    REAL,  
    PRIMARY KEY(bar, beer),  
    FOREIGN KEY(beer)  
        REFERENCES Beers(name)  
        ON DELETE SET NULL  
        ON UPDATE CASCADE  
);
```

Attribute-Based Check

- Constraint on the value of a particular attribute.
- CHECK(<condition>) may be added to the declaration for the attribute.
 - Condition must evaluate to TRUE or UNKNOWN.
 - It's an Error if the check condition evaluates to FALSE.
- The condition may refer to the attribute of the relation that is being checked.
- For the condition to reference any other tuples or relations, a subquery must be used.
 - Note: Many Database Management Systems, including PostgreSQL, do not support CHECK with subquery.

Example: Attribute-Based Check

```
CREATE TABLE Sells (  
    bar    CHAR(20),  
    beer   CHAR(20)  CHECK ( beer IN  
                           (SELECT name FROM Beers) ),  
    price  REAL      CHECK ( price <= 5.00 ),  
    PRIMARY KEY (bar, beer)  
);
```

Example: Named Check Constraints

```
CREATE TABLE Sells (  
    bar CHAR(20),  
    beer CHAR(20) CHECK ( beer IN  
                            (SELECT name FROM Beers) ),  
    price REAL  
    CONSTRAINT price_is_cheap  
        CHECK ( price <= 5.00 ),  
    PRIMARY KEY (bar, beer)  
);
```

```
ALTER TABLE Sells DROP CONSTRAINT price_is_cheap;  
ALTER TABLE Sells ADD CONSTRAINT price_is_cheap  
    CHECK ( price <= 5.00 );
```

Timing of Attribute-Based Checks

Attribute-based checks are performed only when a value for that attribute is inserted or updated.

- **Example:**

CHECK (price \leq 5.00) checks every new price and rejects the modification (for that tuple) if the price in Sells is more than \$5.

- **Example:**

CHECK (beer IN (SELECT name FROM Beers)) is not checked if a beer is deleted from Beers (unlike for Foreign Keys).

Tuple-Based Checks

- CHECK (<condition>) may be added as a relation-schema element.
- The condition may refer to any attributes of the relation (in the same tuple).
- But for the condition to reference any other tuples or relations, a subquery must be used.
 - Condition is checked **only** on INSERT or UPDATE into relation that has the CHECK.
 - Note: Many Database Management Systems, including PostgreSQL, do not support CHECK with subquery.

Example: Tuple-Based Check

Only Joe's Bar can sell beer for more than \$5. That is:

IF price is more than 5.00 THEN bar must be Joe's Bar.

```
CREATE TABLE Sells (  
  bar          CHAR(20),  
  beer         CHAR(20),  
  price        REAL,  
  CHECK ( price <= 5.00 OR bar = 'Joe's Bar' ),  
  PRIMARY KEY (bar, beer)  
);
```

IF p THEN q is logically equivalent to: NOT p OR q

p is “price is more than 5.00” and q is “bar is Joe's Bar”.

When are Constraints Checked and Handled?

[This won't be on Exams]

- In the CONSTRAINT clause, you may specify a constraint to be DEFERRABLE or NOT DEFERRABLE
 - DEFERRABLE INITIALLY DEFERRED or
 - DEFERRABLE INITIALLY IMMEDIATE
- Within a transaction, SET CONSTRAINTS determines whether:
 - SET CONSTRAINTS { ALL | name [, ...] } IMMEDIATE;
 - Deferrable constraints are checked immediately for each SQL statement, or
 - SET CONSTRAINTS { ALL | name [, ...] } DEFERRED;
 - Checking is deferred until the end of the transaction.

Example: Deferred Constraint

[This won't be on Exams]

```
CREATE TABLE Sells (  
    bar    CHAR(20),  
    beer   CHAR(20)  CHECK ( beer IN  
                           (SELECT name FROM Beers) ),  
    price  REAL  
    CONSTRAINT price_is_cheap  
               CHECK ( price <= 5.00 )  
               DEFERRABLE INITIALLY DEFERRED,  
    PRIMARY KEY (bar, beer)  
);  
  
SET CONSTRAINTS price_is_cheap IMMEDIATE;
```

Assertions

- These are database-schema elements, like relations or views.
- Defined by:
 CREATE ASSERTION <name>
 CHECK (<condition>);
- Condition may refer to any relation or attribute in the database schema.
- (**Not implemented** in Relational DBMS because they're too complicated and too expensive!)

Triggers: Motivation

- Assertions are powerful, it's difficult and expensive to implement them ...
 - ... so they're not implemented by DBMS.
- Attribute-based and tuple-based CHECKs are simpler and cheaper, but their power is very limited.
- **Triggers** let the user, perhaps a Database Administrator, (DBA), determine when to check for conditions and decide what to do when those conditions occur.

Event-Condition-Action Rules

- Another name for “trigger” is an *ECA Rule*, or **Event-Condition-Action** Rule.
 - **Event**: Typically a database modification, such as “INSERT ON Sells”
 - **Condition**: Any SQL boolean-valued expression
 - **Action**: Any SQL statements

Preliminary Example: A Trigger

- Instead of using a Foreign Key constraint and rejecting insertions into `Sells(bar, beer, price)` when the beer doesn't appear as a name in the relation `Beers(name, manf)`
- We can create a Trigger that adds a tuple to Beers with that `beer` as Beers.name, and with a NULL manufacturer.

Example: Trigger Definition

```
CREATE TRIGGER BeerTrig
```

```
AFTER INSERT ON Sells
```

The Event

```
REFERENCING NEW ROW AS NewTuple  
FOR EACH ROW
```

```
WHEN (NewTuple.beer NOT IN  
      (SELECT name FROM Beers))
```

The Condition

```
INSERT INTO Beers(name)  
VALUES(NewTuple.beer);
```

The Action

CREATE TRIGGER

- Either:
 CREATE TRIGGER <name>
- Or:
 CREATE OR REPLACE TRIGGER <name>
 - That's useful if there is (or might be) a Trigger with that name, and you want to modify that Trigger.
- CREATE OR REPLACE can also be used for other CREATE statements.

Options: The Event

- **AFTER INSERT** can be **BEFORE INSERT**.
 - Also, can use **INSTEAD OF INSERT**, e.g., if the relation is a view.
 - A clever way to execute modifications of a view is to have Triggers that translate them to appropriate modifications of the base tables.
- **INSERT** can be **DELETE** or **UPDATE**.
 - And **UPDATE** can be **UPDATE OF** a particular attribute.

Options: FOR EACH ROW

- Each trigger is either a “Row-level” trigger or a “Statement-level” trigger.
- FOR EACH ROW indicates that it's a Row-level trigger.
 - Row-level trigger:
Trigger is executed once for each modified tuple.
- Not having FOR EACH ROW indicates that the trigger is a Statement-level trigger.
 - Statement-level trigger:
Trigger is executed just once for the entire SQL statement, no matter how many tuples are modified by that statement.
 - Some systems don't support Statement-level triggers.
 - We don't show examples of Statement-level triggers.

Options: REFERENCING

- INSERT statement has a new tuple (for Row-level) or a new table (for Statement-level).
 - The “table” is the set of inserted tuples.
- DELETE statement has an old tuple or table.
- UPDATE statement has both old and new.
- In the Trigger, we can reference these by:
[NEW OLD] [ROW TABLE] AS <name>
- Some systems support access to old row using keyword OLD, and to new row using keyword NEW, so REFERENCING clause isn't needed.
- [NEW OLD] TABLE is for Statement-level triggers.

Options: The Condition

- Any boolean-valued condition.
- Evaluated on the database as it existed before or after the Event, depending on whether BEFORE or AFTER is specified on the Event.
 - The Event may be INSERT, DELETE, UPDATE or SELECT.
 - Evaluation of the Condition always occurs before any triggered changes are executed, whether BEFORE Event or AFTER Event (or INSTEAD OF Event) is specified.
- Access the new/old row/table through the names in the REFERENCING clause.
 - Or just use NEW and OLD, in systems which support that.

Options: The Action

- There can be more than one SQL statement in the action.
 - Surround statement by BEGIN . . . END if there is more than one statement.
- Queries don't make sense in a triggered action, so statements in Trigger actions should be modification statements.
 - Some DBMS allow Trigger actions to be Stored Procedures, which we'll discuss later in the quarter.
 - PostgreSQL requires that Trigger action be a single Stored Procedure.

Another Example

- Using `Sells(bar, beer, price)` and a unary relation `RipoffBars(bar)`, maintain a list of RipoffBars that raise the price of some beer by more than \$1.

The Trigger

